

# Introduction to R and first steps in categorical data analysis

Natalia Levshina © 2017

# Outline

1. Introduction to R
2. Basics of R syntax
3. Main objects in R
4. Creating and importing your data into R
5. First inspection of data
6. Work with ParTy

# What is R?

- statistical computing environment (from *t*-test to generalized linear models, and more...)
  - core distribution “base”
  - add-on packages (> 10K as of June 2017)
- programming language
- tools for creation of publication-quality plots

# Where to get R?

- Distribution and packages: CRAN (Comprehensive R Archive Network) <http://cran.r-project.org/>
- Information: <http://www.r-project.org/>

# RStudio

- Highly recommended (easy to manage projects, packages, data, graphs, etc.)!
- Available from <http://www.rstudio.com/products/RStudio/>

# Outline

1. Introduction to R
2. Basics of R syntax
3. Main objects in R
4. Creating and importing your data into R
5. First inspection of data
6. Work with ParTy

# Input and output

```
> 2 + 2
```

```
[1] 4
```

```
> sample(100, 25) #random sampling of 25 elements  
from integers 1 to 100
```

```
[1] 49 45 70 51 54 5 7 19 60 82 35 55 6 76 93  
89 44
```

```
[18] 8 48 87 53 34 86 96 63
```

# Basic arithmetic functions

```
> 25^2
```

```
[1] 625
```

```
> 625^0.5
```

```
[1] 25
```

```
> sqrt(625)
```

```
[1] 25
```

```
> log(5)
```

```
[1] 1.609438
```



# Creation of objects

```
> a <- 3
```

```
> a
```

```
[1] 3
```

```
> a + 5
```

```
[1] 8
```

# Exercise

- The population of Finland is approximately 5.5M people, and the population of Burkina Faso is about 17.3M. Create two R numeric vectors with 1 element in each, named Finland and Burkina\_Faso, with the corresponding population values. Compute their sum.

# Beware: = and ==

```
> a = 3 # creates an object a with the value 3, an  
alternative to "a <- 3"
```

```
> a == 3 # tests if a equals 3
```

```
[1] TRUE
```

```
> a == 10 # tests if a equals 10
```

```
[1] FALSE
```

# R is case-sensitive!

```
> b <- 7
```

```
> a + b
```

```
[1] 10
```

```
> a + B
```

```
Error: object 'B' not found
```

# Managing your objects

```
> ls() #returns a list of objects
```

```
[1] "a"      "b"
```

```
> rm(b) #removes an object
```

```
> ls()
```

```
[1] "a"
```

# Saving your workspace

1. Click on the cross or type

```
> q()
```

Select the action (to save or not to save).

```
> getwd() #to find out where your workspace will  
be saved
```

```
[1] "C:/Users/Your/Directory"
```

```
> setwd("C:/Users/Your/Directory") #to change it,  
if you like
```

2. Next session: restart R or, if you have many different workspaces, click on the R from the directory; alternatively:

```
> load("yourDirectory/yourFile.RData")
```

# Getting help

```
> ?cor #to open a help file with information about  
function 'cor'
```

```
> ??correlation #returns a list of functions that  
contain this expression
```

# Exercise

- Get help on the function `summary()`.



# Errors

```
> x <- 1:10 # creates a numeric vector with  
integers from 1 to 10
```

```
> x
```

```
[1]  1  2  3  4  5  6  7  8  9 10
```

```
> meann(x) # we want to compute the mean value of  
x: a typo
```

```
Error: could not find function "meann"
```

```
> mean(x) # correct
```

```
[1] 5.5
```

# Warning messages

```
> mytable <- rbind(c(1, 2), c(3, 4)) #create a 2-  
by-2 table
```

```
> mytable
```

	[,1]	[,2]
[1,]	1	2
[2,]	3	4

```
> chisq.test(mytable)
```

Pearson's Chi-squared test with Yates' continuity  
correction

data: mytable

X-squared = 0, df = 1, p-value = 1

**Warning message:**

**In chisq.test(mytable) : Chi-squared approximation  
may be incorrect**

# Outline

1. Introduction to R
2. Basics of R syntax
3. Main objects in R
4. Creating and importing your data into R
5. First inspection of data
6. Work with ParTy

# Important data types in R

- Numeric vectors
- Character vectors
- Factors
- Data frames
- Contingency tables
- Matrices
- Distance matrices

# Numeric vectors

```
> vnum <- 1:5 # a vector of integers from 1 to 5
> vnum
[1] 1 2 3 4 5
> is(vnum)
[1] "integer"          "numeric"
"vector"
[...]
```

If not a sequence:

```
> RT <- c(455, 773, 512, 667) #reaction times in an
experiment
> RT
[1] 455 773 512 667
```

# Character vectors

```
> sex <- c("f", "m", "m", "f")
```

```
> sex
```

```
[1] "f" "m" "m" "f"
```

```
> is(sex)
```

```
[1] "character"
```

```
"vector"
```

```
[...]
```

# Factors

```
> sex.f <- factor(sex)
```

```
> sex.f
```

```
[1] f m m f
```

```
Levels: f m
```

```
> is(sex.f)
```

```
[1] "factor"
```

```
[...]
```

```
"integer"
```

# Data frames

```
> mydf <- data.frame(sex, RT) #char. vectors turn  
into factors
```

```
> mydf
```

	sex	RT
1	f	455
2	m	773
3	m	512
4	f	667

```
> is(mydf)
```

```
[1] "data.frame" "list" [...]
```



# Exercise

Create a character vector with the names of your fellow students. Create a vector with their heights (in cm). Combine the vectors in one data frame.

# Contingency tables

- Let's add another factor to the dataframe, *dialect*:

```
> mydf$dialect <- c("BrE", "AmE", "AmE", "BrE")
```

```
> mydf
```

	sex	RT	dialect
1	f	455	BrE
2	m	773	AmE
3	m	512	AmE
4	f	667	BrE

```
> table(mydf$sex, mydf$dialect)
```

	AmE	BrE
f	0	2
m	2	0

# Matrices

```
> m <- cbind(1:5, 10:6)
```

```
> m
```

	[,1]	[,2]
[1,]	1	10
[2,]	2	9
[3,]	3	8
[4,]	4	7
[5,]	5	6

```
> is(m)
```

```
[1] "matrix"      "array"      [...]
```

# Distance matrices

> `eurodist`

[output omitted: distances between several European cities]

My journey yesterday:

	Frankfurt	Stockholm	Tampere
Frankfurt	0	1186	1572
Stockholm	1186	0	395
Tampere	1572	395	0

# My journey

```
> mydist <- rbind(Frankfurt = c(0, 1186, 1572),  
Stockholm = c(1186, 0, 395), Tampere = c(1572,  
395, 0))
```

```
> colnames(mydist) <- rownames(mydist)
```

```
> mydist
```

	Frankfurt	Stockholm	Tampere
Frankfurt	0	1186	1572
Stockholm	1186	0	395
Tampere	1572	395	0

```
> is(mydist)
```

```
[1] "matrix"      "array"      "mMatrix"  
"structure" "vector"
```

# From matrix to distance matrix

```
> mydist <- as.dist(mydist)
```

```
> mydist
```

	Frankfurt	Stockholm
Stockholm	1186	
Tampere	1572	395

```
> is(mydist)
```

```
[1] "dist"
```

... and back

```
> m <- as.matrix(mydist)
```

```
> m
```

	Frankfurt	Stockholm	Tampere
Frankfurt	0	1186	1572
Stockholm	1186	0	395
Tampere	1572	395	0

# Exercise

- Make your own distance matrix, depending on where you have travelled from.



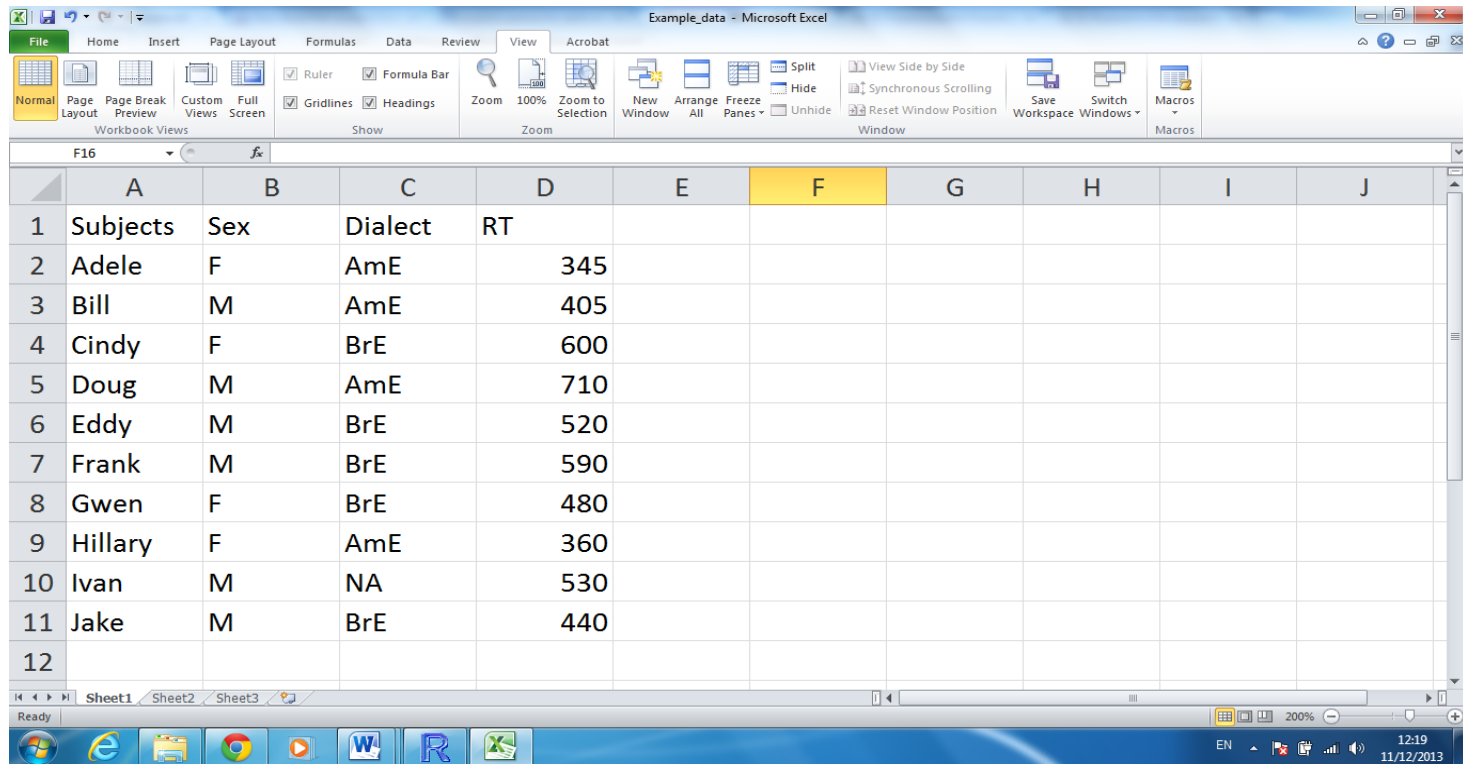
# Quest

1. Compute the square root of 1681.
2. Type in R: `set.seed(x)`, where x is the result of step 1.
3. Create a random sample of 100 numbers from 1 to 100.
4. Find the 20<sup>th</sup> element. This will be your y.
5. Take the y<sup>th</sup> letter in the English alphabet. Write down the letter.
6. Open the help page of the function `read.table` and find the subsection “See also”. Find the first R function mentioned in that subsection. Remove the first letter and write down the result.
7. Find R citation information using `citation()`. Take the 3<sup>rd</sup> word and write down the letter.
8. Put all words together!

# Outline

1. Introduction to R
2. Basics of R syntax
3. Main objects in R
4. Creating and importing your data into R
5. First inspection of data
6. Work with ParTy

# Importing your data to R



The screenshot shows a Microsoft Excel window titled "Example\_data - Microsoft Excel". The ribbon is set to "View", showing options like "Ruler", "Formula Bar", "Gridlines", "Headings", "Zoom", "New Window", "Arrange All", "Freeze Panes", "Split", "Hide", "Unhide", "View Side by Side", "Synchronous Scrolling", "Reset Window Position", "Save Workspace", "Switch Windows", and "Macros". The worksheet contains a table with 12 rows and 11 columns (A-J). The data is as follows:

	A	B	C	D	E	F	G	H	I	J
1	Subjects	Sex	Dialect	RT						
2	Adele	F	AmE	345						
3	Bill	M	AmE	405						
4	Cindy	F	BrE	600						
5	Doug	M	AmE	710						
6	Eddy	M	BrE	520						
7	Frank	M	BrE	590						
8	Gwen	F	BrE	480						
9	Hillary	F	AmE	360						
10	Ivan	M	NA	530						
11	Jake	M	BrE	440						
12										

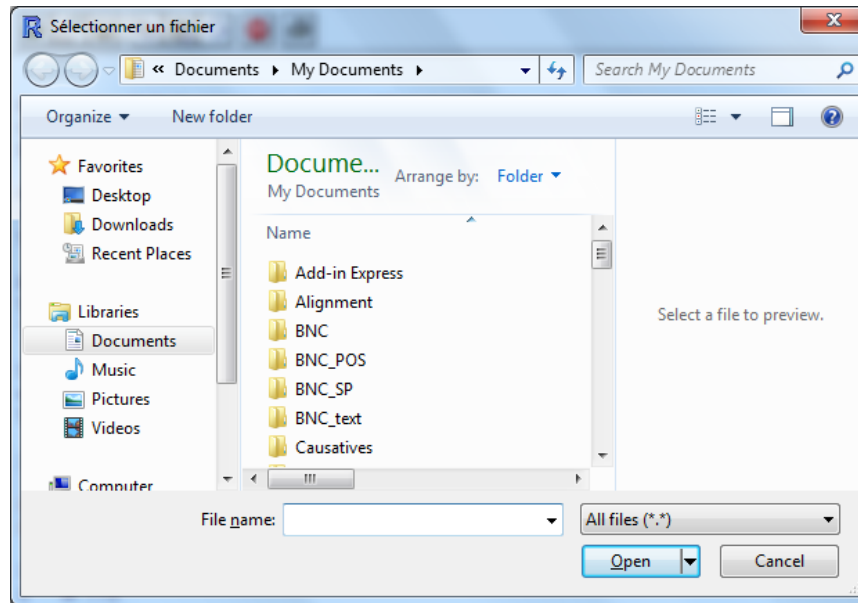
The status bar at the bottom shows "Ready", "Sheet1", "Sheet2", "Sheet3", "200%", and the date/time "12:19 11/12/2013".

# Importing your data into R

1. Create a similar table in Excel (or OpenOffice Calc). Don't forget to create a header. In case of missing values, put NA. No empty cells!
2. Save the file as a tab delimited text file (.txt).
3. Read the file in R:

```
> mydata <- read.table(file = file.choose(), header  
= TRUE)
```

# Interactive choice



# Exercise

Create the following table in Excel (or OpenOffice Calc) and import it in R as a data frame under the name *Linguists*.

Last name	First name	Framework	Born	Died
de Saussure	Ferdinand	Structuralism	1857	1913
Chomsky	Noam	Generative Linguistics	1928	NA
Lakoff	George	Cognitive Linguistics	1941	NA

# Outline

1. Introduction to R
2. Basics of R syntax
3. Main objects in R
4. Creating and importing your data into R
5. First inspection of data
6. Work with ParTy

# Summarizing the data

```
> summary(mydf)
```

sex	RT	dialect
f:2	Min. :455.0	Length:4
m:2	1st Qu.:497.8	Class :character
	Median :589.5	Mode :character
	Mean :601.8	
	3rd Qu.:693.5	
	Max. :773.0	

```
> str(mydf)
```

```
'data.frame': 4 obs. of 3 variables:  
 $ sex      : Factor w/ 2 levels "f","m": 1 2 2 1  
 $ RT       : num  455 773 512 667  
 $ dialect: chr  "BrE" "AmE" "AmE" "BrE"
```



# Selecting observations

```
> mydf[1,]
```

```
  sex  rt  dialect      #the first row  
1   f 455   BrE
```

```
> mydf[,2]
```

```
[1] 455 773 512 667 #the second column
```

```
> mydf[1,2]
```

```
[1] 455 #the element in the first row, second  
column
```

# Using logical operators

```
> mydf[mydf$sex == "f",]
```

	sex	RT	dialect
1	f	455	BrE
4	f	667	BrE

```
> mydf[mydf$sex != "m", ]
```

	sex	RT	dialect
1	f	455	BrE
4	f	667	BrE

```
> mydf[mydf$RT < 500,]
```

	sex	RT	dialect
1	f	455	BrE

# Exercise

- Read the information about the built-in dataset sleep.
- Find how many and which subjects had actually a decrease in sleep.
- Make a subset of all subjects that belong to group 2.

# English Lexicon Project data

```
> str(ELP)
```

```
'data.frame':      880 obs. of  5 variables:
 $ Word      : Factor w/ 880 levels
 "abbreviation",...: 631 747 200 773 821 134 845 140
 94 354 ...
 $ Length    : int   7 10 10 8 6 5 5 8 8 6 ...
 $ SUBTLWF   : num   0.96 4.24 0.04 1.49 1.06 3.33 0.1
 0.06 0.43 5.41 ...
 $ POS       : Factor w/ 3 levels "JJ","NN","VB": 2 2
 3 2 2 2 3 2 2 2 ...
 $ Mean_RT   : num   791 693 960 771 882 ...
```

# 2 ways to tabulate a variable

```
> attach(ELP)
```

```
> summary(POS)
```

```
JJ  NN  VB  
159 532 189
```

```
> table(POS)
```

```
POS  
   JJ   NN   VB  
159 532 189
```

# Proportions and percentages

```
> prop.table(table(POS))
```

```
POS
```

	JJ	NN	VB
	0.1806818	0.6045455	0.2147727

```
> prop.table(table(POS)) * 100
```

```
POS
```

	JJ	NN	VB
	18.06818	60.45455	21.47727

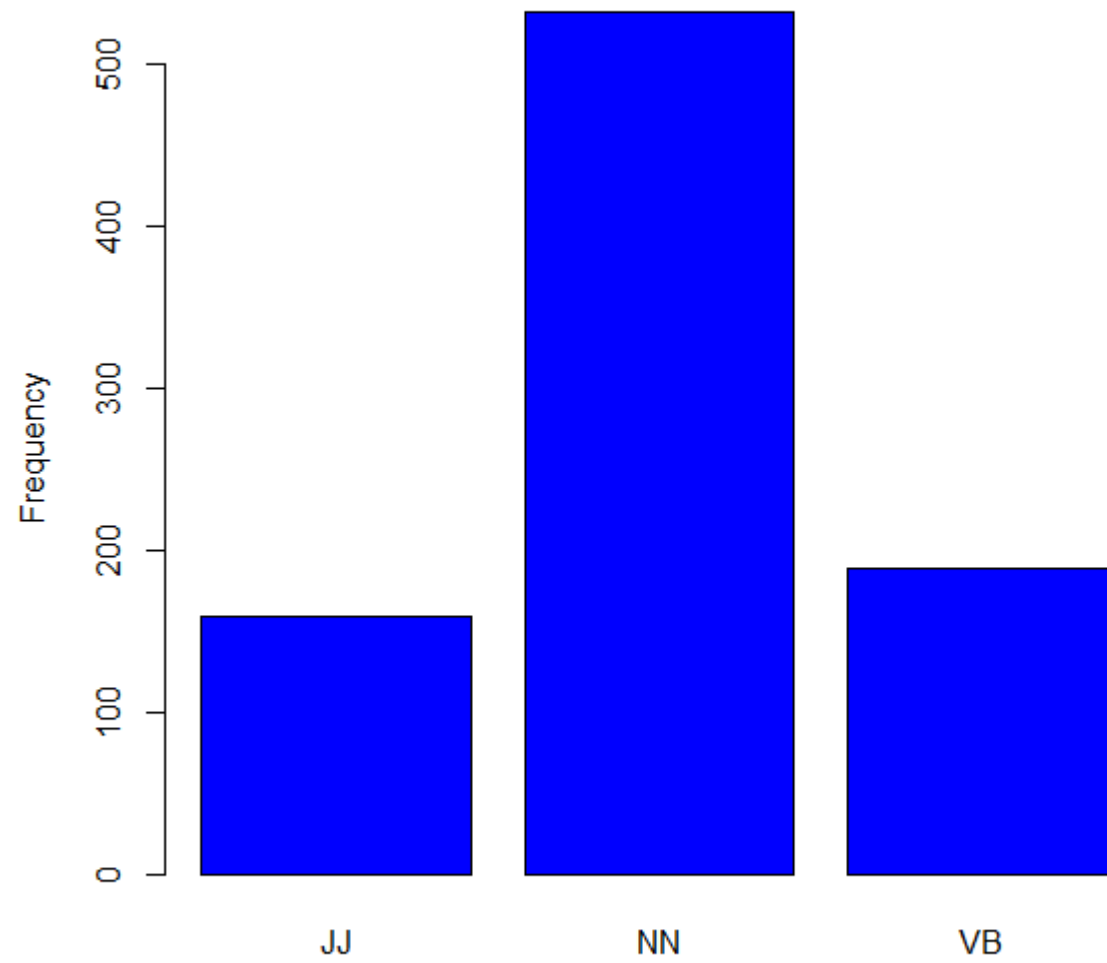
# Bar plot with counts

```
> barplot(table(POS))
```

A fancier version:

```
> barplot(table(POS), col = "blue", main =  
"Frequencies of POS", ylab = "Frequency")
```

**Frequencies of POS**

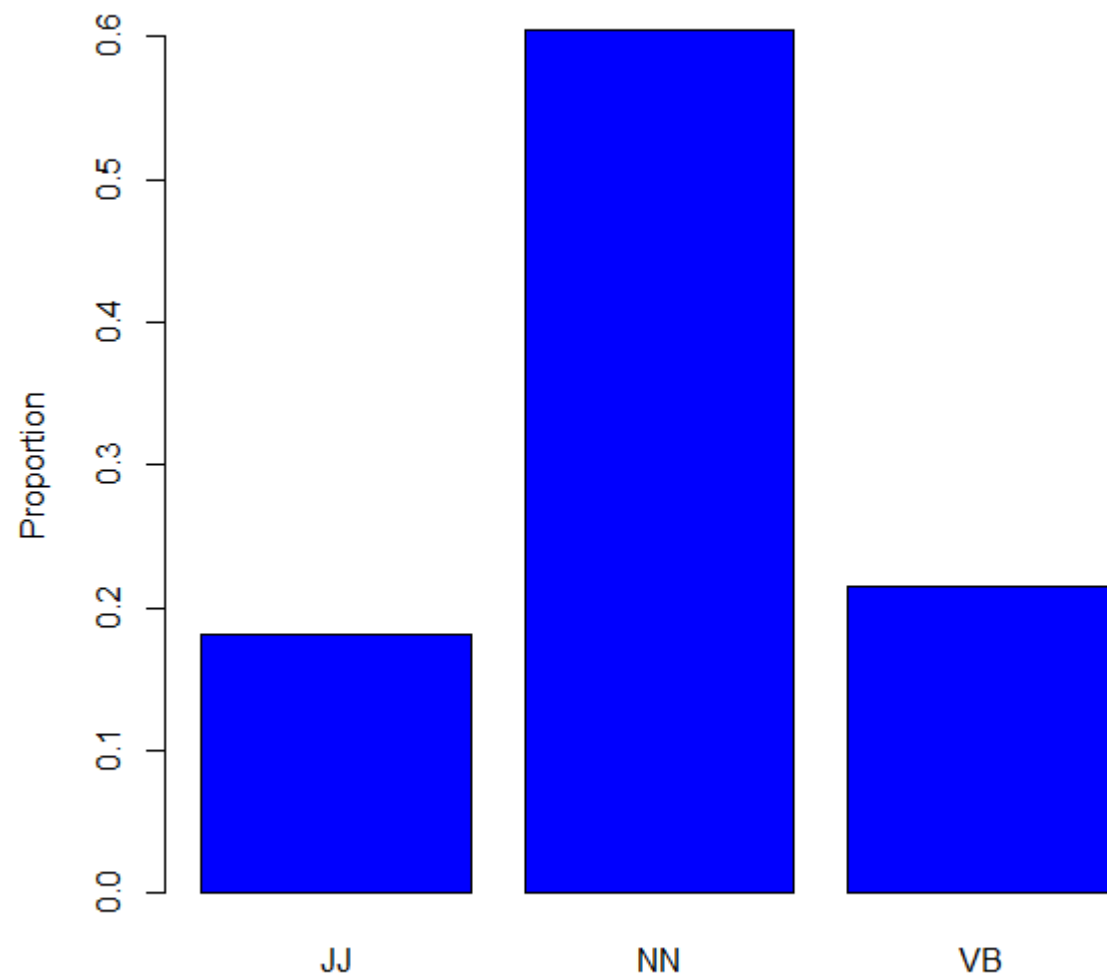




# Bar plot with proportions

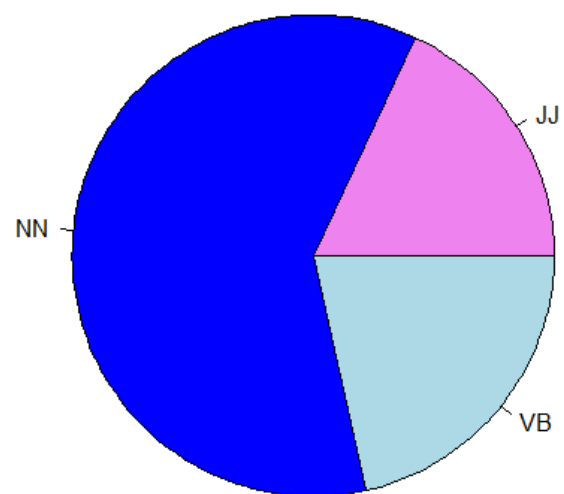
```
> barplot(prop.table(table(POS)), col = "blue",  
main = "Proportions of POS", ylab = "Proportion")
```

**Proportions of POS**



# Pie chart

```
> pie(table(POS), col = c("violet", "blue",  
"lightblue"))
```



# Exercise: Your colours

- Create a factor with your and your fellow students' favourite colours.
- Compute the proportions of each of the colours. Which one is the most popular in the group?
- Create a pie chart with the colours corresponding to each colour category.

# Nerds and geeks

```
> str(nerd)
```

```
'data.frame':      1316 obs. of  5 variables:
 $ Noun       : Factor w/ 2 levels "geek","nerd": 2 2
2 2 2 2 2 2 2 2 ...
 $ Num        : Factor w/ 2 levels "pl","sg": 1 1 1 1
1 1 1 1 1 1 ...
 $ Century    : Factor w/ 2 levels "XX","XXI": 1 2 1
1 1 2 2 1 2 1 ...
 $ Register   : Factor w/ 4 levels
"ACAD","MAG","NEWS",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ Eval       : Factor w/ 3 levels
"Neg","Neutral",...: 2 2 2 2 2 2 2 2 2 2 ..
```

# Noun by Century

```
> attach(nerd)
```

```
> table(Noun, Century)
```

	Century	
Noun	XX	XXI
geek	197	473
nerd	318	328

# Proportions for two-dimensional tables

```
> prop.table(table(Noun, Century)) #all cells sum up to 1
```

```
      Century
Noun      XX      XXI
geek 0.1496960 0.3594225
nerd 0.2416413 0.2492401
```

```
> prop.table(table(Noun, Century), 1) #rows sum up to 1
```

```
      Century
Noun      XX      XXI
geek 0.2940299 0.7059701
nerd 0.4922601 0.5077399
```

```
> prop.table(table(Noun, Century), 2) #columns sum up to 1
```

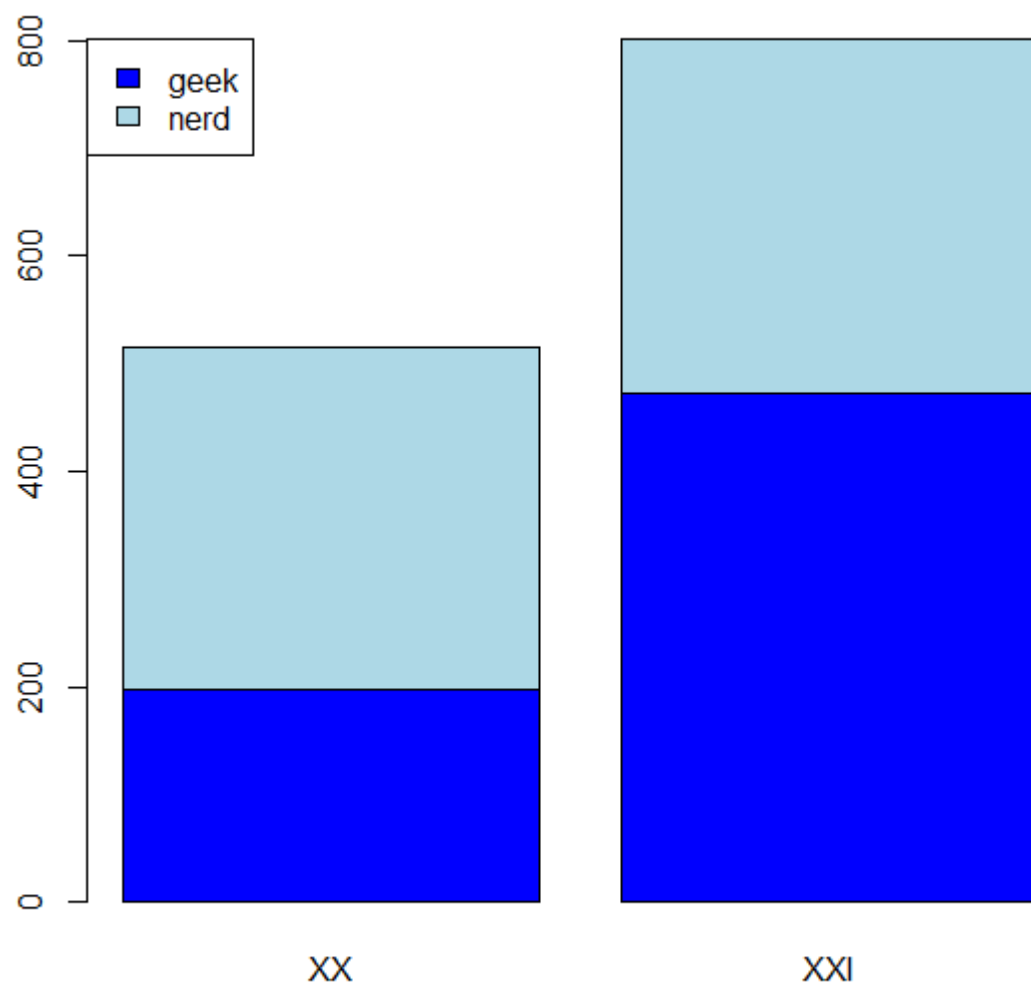
```
      Century
Noun      XX      XXI
geek 0.3825243 0.5905119
nerd 0.6174757 0.4094881
```



# Bar plots with 2 variables

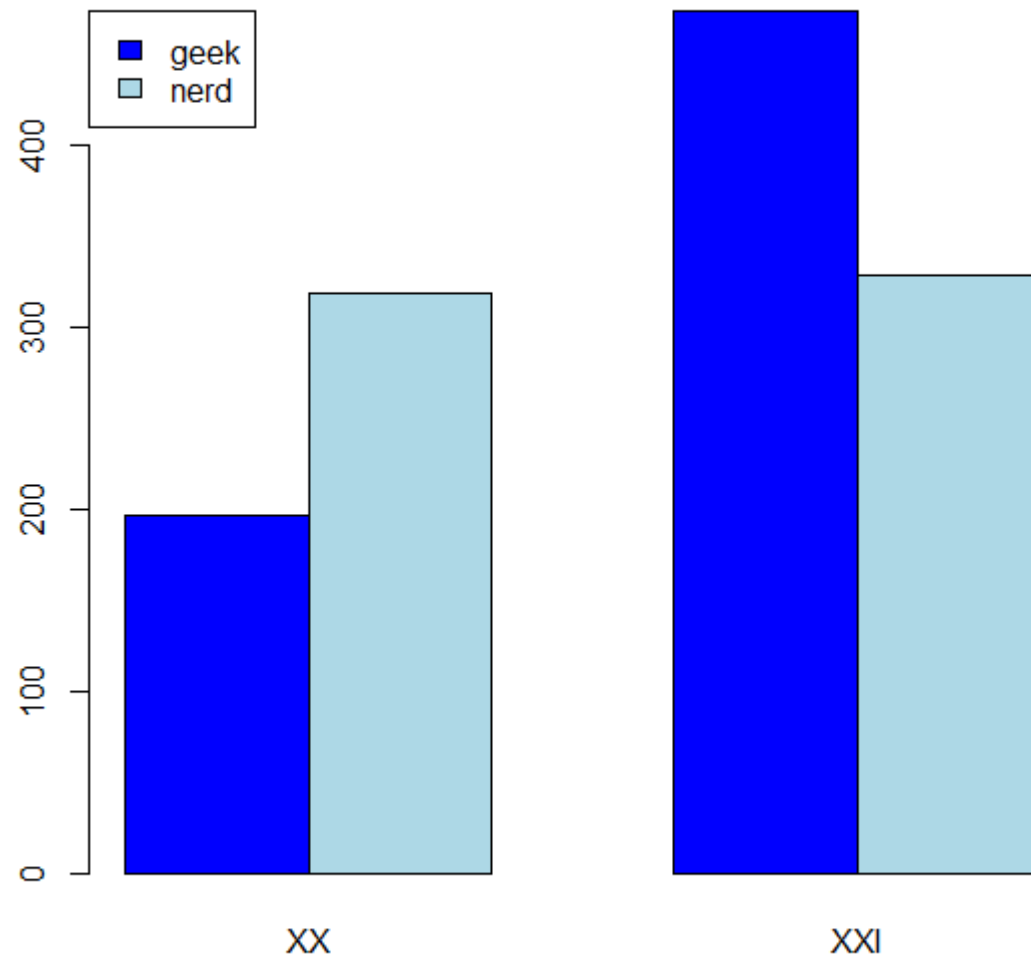
```
> barplot(table(Noun, Century), col = c("blue",  
"lightblue"))
```

```
> legend("topleft", legend = c("geek", "nerd"),  
fill = c("blue", "lightblue"))
```



# Bar plots with unstacked bars

```
> barplot(table(Noun, Century), col = c("blue",  
"lightblue"), beside = TRUE)  
  
> legend("topleft", legend = c("geek", "nerd"),  
fill = c("blue", "lightblue"))
```



# Detach the data frames

```
> detach(ELP)
```

```
> detach(nerd)
```

# Outline

1. Introduction to R
2. Basics of R syntax
3. Main objects in R
4. Creating and importing your data into R
5. Editing your data in R
6. Inspection of categorical data
7. Work with ParTy

# Work with ParTy

1. Choose a linguistic construction in English (e.g. preposition in + NP, imperative).
2. Find the correspondences in different translations. Take at least two languages.
3. Create a dataset in Excel with the original sentences (rows) and translations in the target languages (columns).
4. Import the dataset into R as a data frame. Attach it.
5. Summarize the data in the translations as relative frequencies and percentages.
6. Visualize the data (bar plots, pie charts).
7. Save the plots.